

# Class17 Mini Project

Dennis Kim

## Getting Started

We will start by downloading the most recently dated “Statewide COVID-19 Vaccines Administered by ZIP Code” CSV file from: <https://data.ca.gov/dataset/covid-19-vaccine-progress-dashboard-data-by-zip-code>

```
# Import vaccination data
vax <- read.csv("covid19vaccinesbyzipcode_test.csv")
head(vax)
```

	as_of_date	zip_code	tabulation_area	local_health_jurisdiction	county
1	2021-01-05		95446	Sonoma	Sonoma
2	2021-01-05		96014	Siskiyou	Siskiyou
3	2021-01-05		96087	Shasta	Shasta
4	2021-01-05		96008	Shasta	Shasta
5	2021-01-05		95410	Mendocino	Mendocino
6	2021-01-05		95527	Trinity	Trinity
	vaccine_equity_metric_quartile			vem_source	
1		2	Healthy Places Index Score		
2		2	CDPH-Derived ZCTA Score		
3		2	CDPH-Derived ZCTA Score		
4		NA	No VEM Assigned		
5		3	CDPH-Derived ZCTA Score		
6		2	CDPH-Derived ZCTA Score		
	age12_plus_population	age5_plus_population	tot_population		
1	4840.7		5057		5168
2	135.0		135		135
3	513.9		544		544
4	1125.3		1164		NA
5	926.3		988		997
6	476.6		485		499

	persons_fully_vaccinated	persons_partially_vaccinated
1	NA	NA
2	NA	NA
3	NA	NA
4	NA	NA
5	NA	NA
6	NA	NA

	percent_of_population_fully_vaccinated
1	NA
2	NA
3	NA
4	NA
5	NA
6	NA

	percent_of_population_partially_vaccinated
1	NA
2	NA
3	NA
4	NA
5	NA
6	NA

	percent_of_population_with_1_plus_dose	booster_recip_count
1	NA	NA
2	NA	NA
3	NA	NA
4	NA	NA
5	NA	NA
6	NA	NA

	bivalent_dose_recip_count	eligible_recipient_count
1	NA	0
2	NA	0
3	NA	2
4	NA	2
5	NA	0
6	NA	0

redacted

1	Information redacted in accordance with CA state privacy requirements
2	Information redacted in accordance with CA state privacy requirements
3	Information redacted in accordance with CA state privacy requirements
4	Information redacted in accordance with CA state privacy requirements
5	Information redacted in accordance with CA state privacy requirements
6	Information redacted in accordance with CA state privacy requirements

Q1. What column details the total number of people fully vaccinated?

The column labeled persons fully vaccinated

Q2. What column details the Zip code tabulation area?

The column zip code tabulation area

Q3. What is the earliest date in this dataset?

```
head(sort(vax$as_of_date))
```

```
[1] "2021-01-05" "2021-01-05" "2021-01-05" "2021-01-05" "2021-01-05"  
[6] "2021-01-05"
```

```
#or  
vax$as_of_date[1]
```

```
[1] "2021-01-05"
```

The earliest date in this dataset is 1/5/2021

Q4. What is the latest date in this dataset?

```
head(sort(vax$as_of_date, decreasing=TRUE))
```

```
[1] "2023-02-28" "2023-02-28" "2023-02-28" "2023-02-28" "2023-02-28"  
[6] "2023-02-28"
```

```
#or  
vax$as_of_date[nrow(vax)]
```

```
[1] "2023-02-28"
```

The latest date in this dataset is 2/28/2023

We can use the `skim()` function for a quick overview of a new dataset like this

```
skimr::skim(vax)
```

Table 1: Data summary

Name	vax
Number of rows	199332
Number of columns	18
Column type frequency:	
character	5
numeric	13
Group variables	None

#### Variable type: character

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
as_of_date	0	1	10	10	0	113	0
local_health_jurisdiction	0	1	0	15	565	62	0
county	0	1	0	15	565	59	0
vem_source	0	1	15	26	0	3	0
redacted	0	1	2	69	0	2	0

#### Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
zip_code_tabulation_area	0	1.00	93665.11	1817.38	0	192257.75	3658.50	5380.50	7635.0	
vaccine_equity_metric_0831tile	9831	0.95	2.44	1.11	1	1.00	2.00	3.00	4.0	
age12_plus_population	0	1.00	18895.04	8993.87	0	1346.95	13685.13	1756.18	8556.7	
age5_plus_population	0	1.00	20875.24	1105.97	0	1460.50	15364.08	14877.00	1902.0	
tot_population	9718	0.95	23372.77	2628.51	2	2126.00	18714.08	168.00	11165.0	
persons_fully_vaccinated	16525	0.92	13962.35	5054.09	1	930.00	8566.00	23302.08	7566.0	
persons_partially_vaccinated	16525	0.92	1701.64	2030.18	11	165.00	1196.00	2535.00	39913.0	
percent_of_population_2025_vaccinated	20825	0.90	0.57	0.25	0	0.42	0.60	0.74	1.0	
percent_of_population_2025_fully_vaccinated	20825	0.90	0.08	0.09	0	0.05	0.06	0.08	1.0	
percent_of_population_2025_1_plus_dose	21859	0.89	0.63	0.24	0	0.49	0.67	0.81	1.0	

skim_variable	n_missing	complete	mean	sd	p0	p25	p50	p75	p100	hist
booster_recip_count	72872	0.63	5837.317165.81	11	297.00	2748.00	9438.25	9553.0		
bivalent_dose_recip_count	158664	0.20	2924.933583.45	11	190.00	1418.00	4626.25	7458.0		
eligible_recipient_count	0	1.00	12801.844908.33	0	504.00	6338.00	21973.00	87234.0		

Q5. How many numeric columns are in this dataset?

There are 13 numeric columns

Q6. Note that there are “missing values” in the dataset. How many NA values there in the persons\_fully\_vaccinated column?

```
sum(is.na(vax$persons_fully_vaccinated))
```

```
[1] 16525
```

There are 16,525 NA values in the persons fully vaccinated column

Q7. What percent of persons\_fully\_vaccinated values are missing (to 2 significant figures)?

```
round(((sum(is.na(vax$persons_fully_vaccinated))/nrow(vax))*100), 2)
```

```
[1] 8.29
```

8.29% of the persons\_fully\_vaccinated values are missing

Q8. [Optional]: Why might this data be missing?

Large parts of zip codes are federal land that do not report their data

## Working With Dates

One of the “character” columns of the data is as\_of\_date, which contains dates in the Year-Month-Day format.

Dates and times can be annoying to work with at the best of times. However, in R we have the excellent lubridate package, which can make life allot easier. Here is a quick example to get started.

```
library(lubridate)
```

Attaching package: 'lubridate'

The following objects are masked from 'package:base':

```
date, intersect, setdiff, union
```

What is today's date (at the time I am writing this)

```
today()
```

```
[1] "2023-03-07"
```

The `as_of_date` column of our data is currently not that usable. For example we can't easily do math with it like answering the simple question how many days have passed since data was first recorded.

```
# This will give an Error!  
# today() - vax$as_of_date[1]
```

However if we convert our date data into a lubridate format things like this will be much easier as well as plotting time series data later on.

```
# Specify that we are using the year-month-day format  
vax$as_of_date <- ymd(vax$as_of_date)
```

Math with dates, how many days old am I?

```
today() - ymd("2001-11-21")
```

Time difference of 7776 days

Now we can do math with dates. For example: How many days have passed since the first vaccination reported in this dataset?

```
today() - vax$as_of_date[1]
```

Time difference of 791 days

Using the last and the first date value we can now determine how many days the dataset span.

```
vax$as_of_date[nrow(vax)] - vax$as_of_date[1]
```

Time difference of 784 days

Q9. How many days have passed since the last update of the dataset?

```
today() - vax$as_of_date[nrow(vax)]
```

Time difference of 7 days

7 days have passed

Q10. How many unique dates are in the dataset (i.e. how many different dates are detailed)?

```
length(unique(vax$as_of_date))
```

```
[1] 113
```

113 unique dates

alternate approaches

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
n_distinct(vax$as_of_date)
```

```
[1] 113
```

## Working with ZIP codes

ZIP codes are also rather annoying things to work with as they are numeric but not in the conventional sense of doing math. One of the numeric columns in the dataset (namely `vax$zip_code_tabulation_area`) are actually ZIP codes - a postal code used by the United States Postal Service (USPS). In R we can use the `zipcodeR` package to make working with these codes easier. For example, let's install and then load up this package and to find the centroid of the La Jolla 92037 (i.e. UC San Diego) ZIP code area.

```
library(zipcodeR)
```

```
geocode_zip('92037')
```

```
# A tibble: 1 x 3
  zipcode lat lng
  <chr>   <dbl> <dbl>
1 92037   32.8 -117.
```

Calculate the distance between the centroids of any two ZIP codes in miles, e.g.

```
zip_distance('92037','92703')
```

```
zipcode_a zipcode_b distance
1      92037      92703      77.3
```

More usefully, we can pull census data about ZIP code areas (including median household income etc.). For example:

```
reverse_zipcode(c('92037', "92703"))
```

```
# A tibble: 2 x 24
  zipcode zipcode_~1 major~2 post_~3 common_c~4 county state lat lng timez~5
  <chr>   <chr>       <chr>   <chr>      <blob> <chr>   <chr> <dbl> <dbl> <chr>
1 92037  92037000000  92037  92037  92037  92037  92037  92037  32.8 -117. 0800
2 92703  92703000000  92703  92703  92703  92703  92703  92703  33.7 -117. 0800
```



```

1 92037   Standard   La Jol~ La Jol~ <raw 20 B> San D~ CA      32.8 -117. Pacific
2 92703   Standard   Santa ~ Santa ~ <raw 21 B> Orang~ CA      33.8 -118. Pacific
# ... with 14 more variables: radius_in_miles <dbl>, area_code_list <blob>,
#   population <int>, population_density <dbl>, land_area_in_sqmi <dbl>,
#   water_area_in_sqmi <dbl>, housing_units <int>,
#   occupied_housing_units <int>, median_home_value <int>,
#   median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
#   bounds_north <dbl>, bounds_south <dbl>, and abbreviated variable names
#   1: zipcode_type, 2: major_city, 3: post_office_city, ...

```

Optional: We can use this `reverse_zipcode()` to pull census data later on for any or all ZIP code areas we might be interested in.

```

# Pull data for all ZIP codes in the dataset
#zipdata <- reverse_zipcode( vax$zip_code_tabulation_area )

```

## Focus on the San Diego area

Let's now focus in on the San Diego County area by restricting ourselves first to `vax$county == "San Diego"` entries. We have two main choices on how to do this. The first using base R the second using the `dplyr` package:

```

# Subset to San Diego county only areas
sd <- vax[ vax$county == "San Diego", ]

```

Using `dplyr` the code would look like this:

```

library(dplyr)

sd <- filter(vax, county == "San Diego")

nrow(sd)

```

```
[1] 12091
```

Using `dplyr` is often more convenient when we are subsetting across multiple criteria - for example all San Diego county areas with a population of over 10,000.

```
sd.10 <- filter(vax, county == "San Diego" &
                age5_plus_population > 10000)
nrow(sd.10)
```

[1] 8588

```
nrow(sd.10)
```

[1] 8588

```
n_distinct(sd.10$zip_code_tabulation_area)
```

[1] 76

Q11. How many distinct zip codes are listed for San Diego County?

```
length(unique(sd$zip_code_tabulation_area))
```

[1] 107

There are 107 distinct zip codes for San Diego

Q12. What San Diego County Zip code area has the largest 12 + Population in this dataset?

```
sd$zip_code_tabulation_area[which.max(sd$age12_plus_population)]
```

[1] 92154

92154 has the largest 12+ population in this data set

Q13. What is the overall average “Percent of Population Fully Vaccinated” value for all San Diego “County” as of the most recent date “2023-2-28”?

```
vax$as_of_date[nrow(vax)]
```

[1] "2023-02-28"

```
# using dplyr
sd.today <- filter(sd, as_of_date == "2023-02-28")

mean(sd.today$percent_of_population_fully_vaccinated, na.rm=TRUE)
```


```
[1] 0.7400878
```

74% are fully vaccinated

Q14. Using either ggplot or base R graphics make a summary figure that shows the distribution of Percent of Population Fully Vaccinated values as of “2022-11-15”?

```
hist(sd.today$percent_of_population_fully_vaccinated, breaks=20)
```

**Histogram of sd.today\$percent\_of\_population\_fully\_vaccinated**




```
library(ggplot2)
ggplot(sd.today) +
  aes(x=percent_of_population_fully_vaccinated) +
  geom_histogram() +
  labs (title = "Vaccination Rate Across San Diego as of Last Week",
        x = "Percent of Population Fully Vaxed",
```

```
y = "Count (ZIP code area)"
```

``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.

Warning: Removed 8 rows containing non-finite values (``stat_bin()``).



## Focus on UCSD/La Jolla


UC San Diego resides in the 92037 ZIP code area and is listed with an age 5+ population size of 36,144.

```
ucsd <- filter(sd, zip_code_tabulation_area=="92037")
ucsd[1,]$age5_plus_population
```

```
[1] 36144
```

Q15. Using ggplot make a graph of the vaccination rate time course for the 92037 ZIP code area:

```
ucsdplot <- ggplot(ucsd) +
  aes(x = as_of_date,
      y = percent_of_population_fully_vaccinated) +
  geom_point() +
  geom_line(group=1) +
  ylim(c(0,1)) +
  labs(title = "Vaccination Rate for La Jolla 92109", x= "Date", y="Percent Vaccinated")
ucsdplot
```



## Comparing to similar sized areas


```
# Subset to all CA areas with a population as large as 92037
vax.36 <- filter(vax, age5_plus_population > 36144 &
  as_of_date == "2023-02-28")
meanline <- mean(vax.36$percent_of_population_fully_vaccinated)
meanline
```

```
[1] 0.7213331
```

```
#head(vax.36)
```

Q16. Calculate the mean “Percent of Population Fully Vaccinated” for ZIP code areas with a population as large as 92037 (La Jolla) as\_of\_date “2022-11-15”. Add this as a straight horizontal line to your plot from above with the `geom_hline()` function?

```
ucsdplot +  
  geom_hline(yintercept=meanline, col="red", lty=2)
```



Q17. What is the 6 number summary (Min, 1st Qu., Median, Mean, 3rd Qu., and Max) of the “Percent of Population Fully Vaccinated” values for ZIP code areas with a population as large as 92037 (La Jolla) as\_of\_date “2023-02-28”?

```
summary(vax.36$percent_of_population_fully_vaccinated)
```


Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.3804	0.6457	0.7181	0.7213	0.7907	1.0000

Q18. Using ggplot generate a histogram of this data.

```
ggplot(vax.36) + aes(x=percent_of_population_fully_vaccinated) + geom_histogram() + labs(x=
```

`stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.

Warning: Removed 2 rows containing missing values (`geom\_bar()`).



Q19. Is the 92109 and 92040 ZIP code areas above or below the average value you calculated for all these above?

```
vax %>% filter(as_of_date == "2022-11-15") %>%  
  filter(zip_code_tabulation_area=="92040") %>%  
  select(percent_of_population_fully_vaccinated)
```

```
percent_of_population_fully_vaccinated  
1 0.548849
```

The returned result is less than what was calculated for the average value above.

Q20. Finally make a time course plot of vaccination progress for all areas in the full dataset with a age5\_plus\_population > 36144.

```
vax.36.all <- filter(vax, age5_plus_population > 36144)

ggplot(vax.36.all) +
  aes(as_of_date,
      percent_of_population_fully_vaccinated,
      group=zip_code_tabulation_area) +
  geom_line(alpha=0.2, color="blue") +
  ylim(0,1) +
  labs(x="Date", y="Percent Vaccinated",
       title="Vaccination rate accross California",
       subtitle="Only areas with population above 36k are shown") +
  geom_hline(yintercept = meanline, linetype=2)
```

Warning: Removed 183 rows containing missing values (`geom\_line()`).

